


# **From Neural Networks to parametric modeling**

- **Cesare Alippi**

# What is machine learning

*« Yes, it is a hot topic »*




# What does it mean “to learn”?

- Hastie, Tibshirani, Friedman:
  - “Vast amounts of data are being generated in many fields, and the statisticians’s job is to make sense of it all: to extract important patterns and trends, and to understand “what the data says”. We call this *learning from data*.”
- Mitchell:
  - “The field of machine learning is concerned with the question of how to construct computer programs that automatically improve with experience.”
- Alippi:
  - “The ultimate goal of machine learning is to provide the simplest consistent method able to explain past and future data without requesting an explicitly programmed algorithm.”


# Mitchell formalization of the learning framework

A computer program is said to **learn** from **experience  $E$**  with respect to some class of **tasks  $T$**  and **performance measure  $P$** , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ .



# The regression problem

- Linear regression
- Nonlinear regression
- Feed-forward neural networks




# Linear regression

*«Keep the model simple»*




# Some examples

- linear regression



# Some examples

- linear regression: least mean square algorithm





# Multiple Linear regression (as they taught you)


A set of  $n$  data couples (training set)

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$$

is given

$$x \in \mathbb{R}^d, y \in \mathbb{R}$$

$x$  is column vector



# Multiple linear regression


**Assume** that the unknown function that generates the data is linear and that there is a gaussian uncertainty affecting measurements in an additive way

$$y(x) = \theta_1^o + \theta_2^o x_2 + \cdots \theta_d^o x_d + \eta$$

Canonical form for system model

$$y(x) = x^T \theta^o + \eta$$

$$\theta^o \in \mathbb{R}^d \quad \eta = N(0, \sigma_\eta^2)$$



Optimal parameters, grouped in a column vector, are unknown, as it is the variance of noise

# Multiple linear regression

We know that the unknown system model is linear. Which **family of models** should we consider to best fit the available data?


$$\hat{y}(x) = f(\theta, x) = x^T \theta$$

How to determine the best estimated parameter  $\hat{\theta}$  so that we generate the model

$$f(\hat{\theta}, x) = x^T \hat{\theta}$$

approximating unknown function  $x^T \theta^o$  ?

# Multiple linear regression



Idea: select the linear function that minimizes the average distance between given points and the linear function: we obtain the **Least Mean Square –LMS-** procedure

# Multiple linear regression

Performance function

$$V_n(\theta) = \frac{1}{n} \sum_{i=1}^n (y(x_i) - f(\theta, x_i))^2$$

The parameter vector to be chosen is


$$\hat{\theta} = \arg \min_{\theta \in \Theta} V_n(\theta)$$

# Linearity must be intended according to the parameters!!!

- Linear regression

$$y = a + bz + cz^2 + dz^3$$

$$\theta = [a, b, c, d]; x = [1, z, z^2, z^3]$$



# Multiple linear regression: how to estimate parameters

By grouping data as

$$Y = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \dots \\ y_n \end{bmatrix} \quad X = \begin{bmatrix} \mathbf{x}_1^T \\ \mathbf{x}_2^T \\ \mathbf{x}_3^T \\ \dots \\ \mathbf{x}_n^T \end{bmatrix}$$

We can rewrite the loss function in a canonical form

$$V_n^*(\theta) = \sum_{i=1}^n (y(x_i) - x_i^T \theta)^2 = (Y - X\theta)^T (Y - X\theta)$$

# Multiple linear regression

Stationary points are those for which

$$\frac{\partial V_n^*(\theta)}{\partial \theta} = -2X^T Y + 2X^T X \theta = 0$$

Therefore, the parameter vector minimizing the performance function is

$$\hat{\theta} = (X^T X)^{-1} X^T Y$$

and the best approximating model is

$$f(\hat{\theta}, x) = x^T \hat{\theta}$$



# Multiple linear regression

- A computer program is said to **learn** from experience  $E$  with respect to some class of tasks  $T$  and performance measure  $P$ , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ .
- Let's list the different actors
  - Task  $T$ : regression
  - Experience  $E$ :  $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$
  - Performance measure  $P$ : Mean square error
  - Performance at task:
$$V_n(\theta) = \frac{1}{n} \sum_{i=1}^n (y(x_i) - f(\theta, x_i))^2$$

# Performance at task: mmh....

- In a typical classroom the teacher provides solution examples/instances related to a concept, e.g., what a cat is.



- However, at exam time, the problems the teacher provides you to test your understanding are not identical to the ones you saw during the course. E.g.,



# Performance at target: mmh....

- Instead, we propose different instances



- Why?
- Performance at task assessed according to

$$V_n(\theta) = \frac{1}{n} \sum_{i=1}^n (y(x_i) - f(\theta, x_i))^2$$

is biased to the training set

# How to measure performance at task?

- Much more to come later in the course
- For now, based on our exam experience, we consider another data set, called **test set**,

$$\{(\bar{x}_1, \bar{y}_1), (\bar{x}_2, \bar{y}_2), \dots, (\bar{x}_l, \bar{y}_l)\}$$


- and evaluate performance on it

$$V_l(\hat{\theta}) = \frac{1}{l} \sum_{i=1}^l (\bar{y}_i - \bar{x}^T \hat{\theta})^2$$

The procedure is named cross-validation

# Nonlinear regression

***«I like straight lines but sometimes curves are appreciable»***



# Non-linear regression: formalization

The *stationary* process generating the data

$$y = g(x) + \eta$$

provides, given input  $x_i$  output

$$y_i = g(x_i) + \eta_i$$

Collect a set of couples  
(training set)

$$Z_N = \{(x_1, y_1), \dots, (x_N, y_N)\}$$

The goal of supervised learning is to design the simplest approximating model able to explain past  $Z_N$  data and future instances provided by the data generating process.

# Non-linear regression: formalization

Model unknown function  $g(x)$

with parameterized family of models  $f(\theta, x)$

Consider loss function

$$L(y(x), f(\theta, x))$$

e.g.,

$$L(y(x), f(\theta, x)) = (y(x) - f(\theta, x))^2$$

# The traditional statistical approach

The structural risk

$$\bar{V}(\theta) = \int L(y, f(\theta, x)) p_{x,y} dx dy$$

$$\theta^o = \arg \min_{\theta \in \Theta} \bar{V}(\theta)$$

The empirical risk

$$V_N(\theta) = \frac{1}{N} \sum_{i=1}^N L(y_i, f(\theta, x_i))$$

$$\hat{\theta} = \arg \min_{\theta \in \Theta} V_N(\theta)$$



# Non-linear regression: formalization

Determine the parameter estimate through a minimization problem

$$\hat{\theta} = \arg \min_{\theta \in \Theta} V_N(\theta)$$

carried out by a learning procedure

$$\theta_{i+1} = \theta_i - \varepsilon_L \frac{\partial V_N(\theta)}{\partial \theta} \big|_{\theta_i}$$

and get model

$$f(\hat{\theta}, x)$$

# Comments

The risk associated with the model  
(generalization ability, performance at task)  
can be decomposed in three terms

$$\bar{V}(\hat{\theta}) = (\bar{V}(\hat{\theta}) - \bar{V}(\theta^o)) + (\bar{V}(\theta^o) - V_I) + V_I$$

# Comments

The **inherent risk**

$$V_I$$

depends only on the structure of the learning problem. This term can be improved only by improving the problem itself (e.g., by reducing the instrumentation noise)

# Comments

The **approximation risk**

$$\bar{V}(\theta^o) - V_I$$

depends only on how close the approximating model family is to the process generating the data. As such, a more appropriate model family reduces the risk

# Comments

The **estimation risk**


$$\bar{V}(\hat{\theta}) - \bar{V}(\theta^o)$$

depends on the effectiveness of the learning procedure. As such, a more effective learning algorithm reduces the risk

All consistent learning procedures are equally effective


# Feed-forward neural networks

*«An interesting computational paradigm»*



# Feedforward Neural Networks

Not rarely  $f(\theta, x)$  is chosen as




# Universal approximation theorem

A feedforward network with a single hidden layer containing a finite number of neurons and a linear output neuron approximates any continuous function defined on compact subsets

K. Hornik, "Approximation Capabilities of Multilayer Feedforward Networks", *Neural Networks*, No.4 Vol. 2, 251–257, 1991



# Approximation performance vs. Model complexity



# Controlling the model complexity

- Trial & error
- Early stopping
- Dropout
- Tikhonov regularization